

Exercise 2 - Random Forest Regression Algorithm

Clara Pichler
Hannah Knapp
Sibel Toprakkiran

Group 18

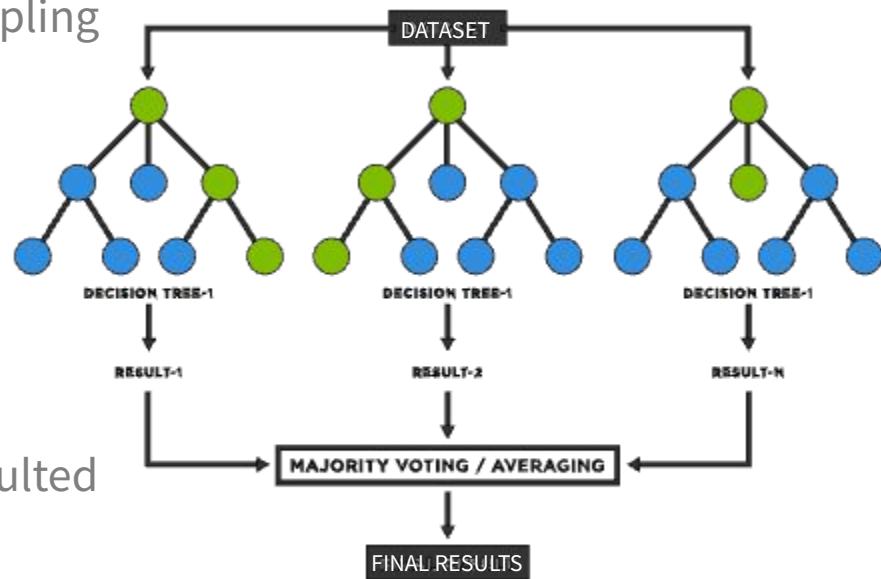
Random Forest Regressor

Step 1: Creating multiple data frames by resampling the original data

-> Bootstrap

Step 2: Build multiple decision trees regressors

Step 3: Averaging over the predicted values resulted from the decision trees



Setup

- Python 3.12.5 - Jupiter Notebook files
- Using only numpy and pandas for the implementation of our regressor
- Testing our model with two dataframes:
 - Airfoil Noise Data
 - Abalone Data
- Using Git to collaborate and merge our code
- Multiple meetings on discord to discuss

Bootstrap

- Method to create multiple subsets of the original data by sampling **with replacement**
- Each subset is used to train one decision tree
- introduces randomness => trees see different views on data
- helps reduce overfitting

```
make_bootstrap(X, y, n_bootstrap = 100)
```

- sample size is the same as the original data
- selecting random data points (indices) from the original data with replacement
- returns a list of tuple containing the bootstrapped X and y

Decision Tree Regressor

Helper Functions:

- Split the data
- Find the best split
- Calculating metrics (for finding the best split)

Tree Building:

- Only using a random subset of the features (`max_features`)
- Recursively splitting data into smaller groups (trees) until a stopping condition is met

Build Decision Tree

```
build_tree(X, y, max_depth, min_samples_split, max_features="sqrt", depth=0,
metric="mse")
```

- building a decision tree recursively
- check if stopping conditions are met (leaf node)
 - `max_depth <= depth` or `len(y) < min_samples_split` or `mse(y) == 0`
 - return mean of `y` if conditions are met
- `feature_subset` random m features without replacement
 - where m is chosen through `max_features`
- `find_best_split` to get `threshold` and `feature_idx`
- `split_dataset` based on this `threshold` and `feature`
- calling `build_tree` for right and left tree with `depth+=1`
- returns dictionary for each tree containing `feature_idx`, `threshold` and left and right tree

Different Metrics to find Best Split

- MSE

- Measures average squared difference of actual and predicted values
- Using weighted version
- We want to minimize weighted MSE (choose lowest)

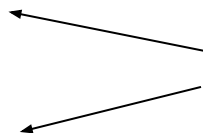
- Variance Reduction

- Measures decrease in variance after splitting
- effective when capturing variability n target values important
- Choosing split with highest variance reduction

- MAE

- Measures average absolute differences of actual and predicted values
- Using weighted version
- minimize weighted MAE

dealing better with
outliers than MSE



Find Best Split

```
find_best_split(X, y, feature_subset=None, metric="mse")
```

- depending on the metric:
 - define `best_metric` as `inf` or `-inf`
 - define `is_better` function
- for each feature in `feature_subset` (if `None` then all features are considered)
 - each value is considered as a threshold
 - based on the splitting the data
 - if one split has length of zero => continue
 - else calculating the metric value
 - with `is_better` function check if we found a better split
 - update `best_metric`, `best_feature` and `best_threshold` if that's the case

RF Regressor Implementation

Class

- `__init__(self, n_trees=10, max_depth=5, min_sample_split=2, max_features="sqrt", metric="mse")`
 - initialising the parameters
- `fit(self, X, y)`
 - creates `n_trees` bootstrap samples with `make_bootstraps`
 - builds a decision tree with each sample using `build_tree`
 - saving them as list of dictionaries in `self.trees`

RF Regressor Implementation - Predicting

```
predict_tree(tree, X)
```

- Traverses the tree recursively, starting at root node and moving down to a leaf node
- If tree is a leaf node
 - returns mean of the target values of the leaf node
- At each non-leaf node decides which child branch to follow
 - `X[feature_idx] <= threshold`, the function proceeds to the left subtree (`tree["left"]`).

```
predict(self, X)
```

- Creates array where each row corresponds to predictions from one tree, and each column corresponds to a single data point
- returns array where each value is the averaged prediction for the corresponding data point

LLM Random Forest Regressor

- Using ChatGPT (4o mini) as our LLM

a very general prompt without giving any details e.g. about hyperparameters

Prompt:

Can you give me the code for a random forest algorithm from scratch with just numpy and pandas? It should be based on Regression Decision Trees.

Comparison with LLM RF Regressor I

logic for building the decision tree and making predictions similar (bootstrapping, building the tree, splitting, finding the best features with threshold)

specific implementation details differ, differences in hyperparameters and modularity

differences in hyperparameters and their values

	Group 18	ChatGPT
Metric (for finding best split)	MSE, MAE, variance reduction	MSE
Feature Subset	can handle sqrt, log, none	none

Comparison with LLM RF Regressor II

differences in fitting phase

- `build_tree` function to create each decision tree as dictionary vs ChatGpt via `DecisionTreeRegressor` class instance.

differences in prediction phase

- iterate over the trees and call `predict_tree` vs. ChatGpt implementation calls the `predict` method of each `DecisionTreeRegressor` class instance.

Comparing with KNN

both can capture non linear relationships

Random Forest

Advantages

- large datasets and high dimensions
- less sensitive to noise and irrelevant features (combining multiple trees)
- can provide information about feature importance
- can be parallelized

Disadvantages

- could be overfitted
- predictions less interpretable
- more complex than KNN, more hyperparameters to be tuned

KNN

Advantages

- easier to understand
- easy to interpret based on nearest neighbours

Disadvantages

- computationally expensive for large datasets
- sensitive to noise

Evaluation

- using two data sets → Airfoil and Abalone
- using different metrics → MSE, MAE and variance reduction
- using different max_features → sqrt(n), log2(n), all features
- running different models with different parameter sets:
 - small n_trees vs. big n_trees
 - small max_depth vs. big n_trees
 - small min_sample split vs. big min_sample split
 - through GridSearchCV, RandomizedSearchCV find best parameters

Default parameters: `n_trees=10, max_depth=5, min_sample_split=2, max_features="sqrt", metric="mse"`

Experimental results I (Airfoil)

Results with variations in metrics

1. mse:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	16.74	13.74	13.22
MAE	3.34	2.88	2.78
R^2	0.65	0.71	0.72

2. variance reduction:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	16.07	-	-
MAE	3.26	-	-
R^2	0.66	-	-

3. mae:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	17.94	-	13.38
MAE	3.40	-	2.78
R^2	0.62	-	0.72

Experimental results I (Abalone)

Results with variations in metrics

1. mse:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.62	5.03	4.91
MAE	1.66	1.56	1.56
R^2	0.49	0.54	0.55

2. variance reduction:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.51	-	-
MAE	1.65	-	-
R^2	0.50	-	-

3. mae:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.70	-	5.22
MAE	1.68	-	1.50
R^2	0.48	-	0.53

Experimental results II (Airfoil)

Results with variations in max_features

1. $\text{max_features} = \sqrt{n}$:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	16.74	-	15.32
MAE	3.34	-	3.12
R^2	0.65	-	0.68

2. $\text{max_features} = \log_2(n)$:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	15.37	-	16.68
MAE	3.10	-	3.29
R^2	0.68	-	0.65

3. $\text{max_features} = n$

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	14.35	13.74	13.92
MAE	2.90	2.88	2.79
R^2	0.70	0.71	0.71

Experimental results II (Abalone)

Results with variations in max_features

1. $\text{max_features} = \sqrt{n}$:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.62	-	5.55
MAE	1.66	-	1.66
R ²	0.49	-	0.50

2. $\text{max_features} = \log_2(n)$:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.84	-	5.78
MAE	1.70	-	1.69
R ²	0.47	-	0.48

3. $\text{max_features} = n$

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	4.89	5.03	5.01
MAE	1.54	1.56	1.57
R ²	0.56	0.54	0.55

Experimental results III (Airfoil)

Results with variations in n_trees

1. n_trees = 10:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	16.74	13.74	13.59
MAE	3.34	2.88	2.84
R^2	0.65	0.71	0.71

2. n_trees = 25:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	15.25	13.13	13.15
MAE	3.15	2.78	2.80
R^2	0.68	0.72	0.72

3. n_trees = 50:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	15.13	13.11	13.03
MAE	3.13	2.79	2.79
R^2	0.68	0.73	0.73

Experimental results III (Abalone)

Results with variations in n_trees

1. n_trees = 10:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.62	5.03	5.12
MAE	1.66	1.56	1.57
R^2	0.49	0.54	0.54

2. n_trees = 25:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.59	4.99	4.99
MAE	1.65	1.56	1.56
R^2	0.49	0.55	0.55

3. n_trees = 50:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.48	4.86	4.92
MAE	1.63	1.54	1.55
R^2	0.50	0.56	0.55

Experimental results IV (Airfoil)

Results with variations in max_depth

1. max_depth = 3:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	24.32	23.28	23.22
MAE	4.02	3.85	3.79
R^2	0.49	0.51	0.51

2. max_depth = 5:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	16.74	13.74	12.81
MAE	3.34	2.88	2.75
R^2	0.65	0.71	0.73

3. max_depth = 7:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	10.19	8.10	7.90
MAE	2.49	2.18	2.17
R^2	0.79	0.83	0.83

Experimental results IV (Abalone)

Results with variations in max_depth

1. max_depth = 3:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	6.58	6.11	6.18
MAE	1.83	1.75	1.75
R^2	0.40	0.45	0.44

2. max_depth = 5:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.62	5.03	5.12
MAE	1.66	1.56	1.57
R^2	0.49	0.54	0.54

3. max_depth = 7:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.14	4.86	4.69
MAE	1.57	1.52	1.51
R^2	0.53	0.56	0.57

Experimental results V (Airfoil)

Results with variations in min_sample split

1. min_sample split = 2:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	16.74	13.74	12.81
MAE	3.34	2.88	2.75
R^2	0.65	0.71	0.73

2. min_sample split = 4:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	16.22	12.76	13.74
MAE	3.26	2.74	2.91
R^2	0.66	0.73	0.71

3. min_sample split = 6:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	17.85	12.41	13.71
MAE	3.39	2.73	2.83
R^2	0.63	0.74	0.71

Experimental results V (Abalone)

Results with variations in min_sample split

1. min_sample split = 2:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.62	5.03	5.12
MAE	1.66	1.56	1.57
R^2	0.49	0.54	0.54

2. min_sample split = 4:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.50	4.91	4.94
MAE	1.63	1.55	1.55
R^2	0.50	0.55	0.55

3. min_sample split = 6:

	RF_18	RF_ChatGPT	RF_Sklearn
MSE	5.52	5.08	4.91
MAE	1.66	1.57	1.54
R^2	0.50	0.54	0.55

Experimental results VI

1. Results with RandomizedSearchCV (Airfoil)

Best Parameters: RF: max_depth=15, max_features=None, n_estimators=150 ; K-NN: n_neighbors=7

	RF_18	RF_ChatGPT	RF_Sklearn	K-NN
MSE	3.77	3.84	3.43	37.86
MAE	1.37	1.38	1.35	4.86
R^2	0.92	0.92	0.93	0.19

2. Results with RandomizedSearchCV (Abalone)

Best Parameters: RF: max_depth=15, max_features='sqrt', min_samples_split=3 ; K-NN: n_neighbors=7, weights='distance'

	RF_18	RF_ChatGPT	RF_Sklearn	K-NN
MSE	4.82	4.78	4.64	4.86
MAE	1.54	1.54	1.51	1.52
R^2	0.56	0.57	0.58	0.54

Conclusions

- Our Random Forest does relatively well in comparison with the RF from ChatGPT and the RF from SKlearn
- There could be improvement in optimization to make the Code faster
- Make the Code modular with adding a Class for the Decision Trees
- Our RF is slower than the RF from Sklearn but faster than the one from ChatGPT
- We get the best results if the number of Trees gets bigger and the max_depth gets bigger
- But to reduce Overfitting the number of max_depth should be lower
- K-NN does not perform as well as RF and for Airfoil it performs really bad
- The Regressors work better for the Airfoil Dataset than the Abalone Dataset